



OLLSCOIL NA GAILLIMHÉ  
UNIVERSITY OF GALWAY

# Ethics in NLP

Dr. Paul Buitelaar  
Data Science Institute, University of Galway



# Learning Objectives

- Understand ethical issues in NLP
- Understand some of the main approaches to address ethical issues in NLP through regulation and technology



## REGULATIONS



**REGULATION (EU) 2016/679 OF THE EUROPEAN PARLIAMENT AND OF THE COUNCIL  
of 27 April 2016  
on the protection of natural persons with regard to the processing of personal data and on the free  
movement of such data, and repealing Directive 95/46/EC (General Data Protection Regulation)**



Support The Guardian

Available for everyone, funded by readers

Contribute →

Subscribe →

Search jobs

Sign in

Search

International  
edition

The  
Guardian

News

Opinion

Sport

Culture

Lifestyle

More

Coronavirus World

BBC

Sign in

Home

News

Sport

Reel

Worklife

Tech Weekly  
Artificial  
intelligence (AI)

NEWS

Home | US Election | Coronavirus | Video | World | UK | Business | Tech | Science | Stories | Enterta

Tech



World

Business

Markets

Breakingviews

Video

More

# 'Dangerous news

By Jane Wakefield  
Technology reporter

27 August 2019

RETAIL OCTOBER 11, 2018 / 12:04 AM / UPDATED 2 YEARS AGO

## Amazon scraps secret AI recruiting tool that showed bias against women

By Jeffrey Dastin

8 MIN READ



## Data centres now use more electricity than run

Electricity used by data centres rises

🕒 Tue, May 3, 2022, 13:30

Updated: Tue, Ma



About

Resources

Subscribe

HOME

COVID-19

FEATURES ▾

SECTORS ▾

APPLICATIONS ▾

TECHNOLOGIES ▾

April 22, 2022

### Unstructured Data Vs. Environment Reflected in Earth Day Commitments

Alex Woodie



Today is Earth Day, which means millions of people are taking time to think about how their actions are impacting the planet. The reflections are especially important for those in the IT business, where every byte of data processed expands our collective carbon footprint. However, data of the unstructured type seems to bear a heavier burden on the earth, especially when it's used to power AI initiatives.

# NLP & Climate Change

# What can be done?



# Regulation



# Regulate the use of AI

RETAIL OCTOBER 11, 2018 / 12:04 AM / UPDATED 2 YEARS AGO

## Amazon scraps secret AI recruiting tool that showed bias against women

By Jeffrey Dastin

Listen Newsletters Shows ▾

MARKETPLACE®

f t i y

DONATE

Q ≡

30 10


New NYC law restricts hiring based on artificial intelligence  
Marketplace


MORE FROM THIS SHOW

↓

⏮ ⏪ ⏩ ⏭


**New NYC law restricts hiring based on artificial intelligence**

 Meghan McCarty Carino | Dec 10, 2021

Heard on:  MARKETPLACE®

▶ LISTEN NOW

◀ ⌕ ⬇



New York City follows Maryland and Illinois in implementing laws aimed at addressing algorithmic discrimination in hiring. Spencer Platt via Getty Images

“When a new law in New York City takes effect at the start of 2023, employers won’t be allowed to use artificial intelligence to screen job candidates **unless the tech has gone through an audit to check for bias.**” – Marketplace, Dec 2021



# Regulation on AI Algorithms



EN English

Search

Shaping Europe's digital future

[Home](#) [Policies](#) [Activities](#) [News](#) [Library](#) [Funding](#) [Calendar](#) [Consultations](#)

[Home](#) > [Policies](#) > Regulatory framework proposal on artificial intelligence

## Regulatory framework proposal on artificial intelligence

The Commission is proposing the first-ever legal framework on AI, which addresses the risks of AI and positions Europe to play a leading role globally.

The regulatory proposal aims to provide AI developers, deployers and users with clear requirements and obligations regarding specific uses of AI. At the same time, the proposal seeks to reduce administrative and financial burdens for business, in particular small and medium-sized enterprises (SMEs).



“The proposed rules will:

- address **risks** specifically created by AI applications;
- propose a list of **high-risk applications**;
- set clear **requirements** for AI systems for high risk applications;
- define specific **obligations** for AI users and providers of high risk applications;
- propose a **conformity assessment** before the AI system is put into service or placed on the market;
- propose **enforcement** after such an AI system is placed in the market;
- propose a **governance structure** at European and national level.”



OLLSCOIL NA GAILLIMHE  
UNIVERSITY OF GALWAY

<https://digital-strategy.ec.europa.eu/en/policies/regulatory-framework-ai>

# Regulation on Data Privacy



# General Data Protection Regulation (GDPR)

“Regulation ... 2016/679 of the European Parliament and of the Council ... regulates the processing by an **individual, a company or an organisation** of **personal data** relating to **individuals** in the EU.”



# Personal Data

- “Personal data is **any information that relates to an identified or identifiable living individual**.
- Different pieces of information, which collected together can lead to the **identification of a particular person**, also constitute personal data.
- Personal data that has been **de-identified, encrypted or pseudonymised but can be used to re-identify a person** remains personal data and falls within the scope of the GDPR.
- **Personal data that has been rendered anonymous** in such a way that the individual is not or no longer identifiable **is no longer considered personal data**.
- For data to be truly anonymised, the **anonymisation must be irreversible.**”



# Personal Data in NLP

**Language data is also personal data** as individuals can be identified by their language data use i.e., how they write or speak

**Natural language processing can lead to ‘fingerprinting’ of individuals**

# Data Protection Impact Assessment

Identify **potential data privacy and data protection issues associated with the data collection for a specific (NLP) task** by answering questions such as:

- *Have you identified the **minimally sufficient amount of data** for your purpose?*
- *Have you **informed affected individuals** what you are doing with their data?*
- *Have you **identified a secondary use for personal data** you collected?*
- *Have you established the **maximum time period required for keeping the data**?*
- *Have you developed **appropriate data anonymization strategies**?*

**Data Protection Impact Assessment** is to be discussed with and approved by the **Data Protection Officer** at the **Data Controller** (organization where the task is done).

# Technology




# Explainable AI

☰ The New York Times Magazine ☷

FEATURE



## Can A.I. Be Taught to Explain Itself?

As machine learning becomes more powerful, the field's researchers increasingly find themselves unable to account for what their algorithms know — or how they know it.





# XAI Concept by DARPA



# XAI Methods

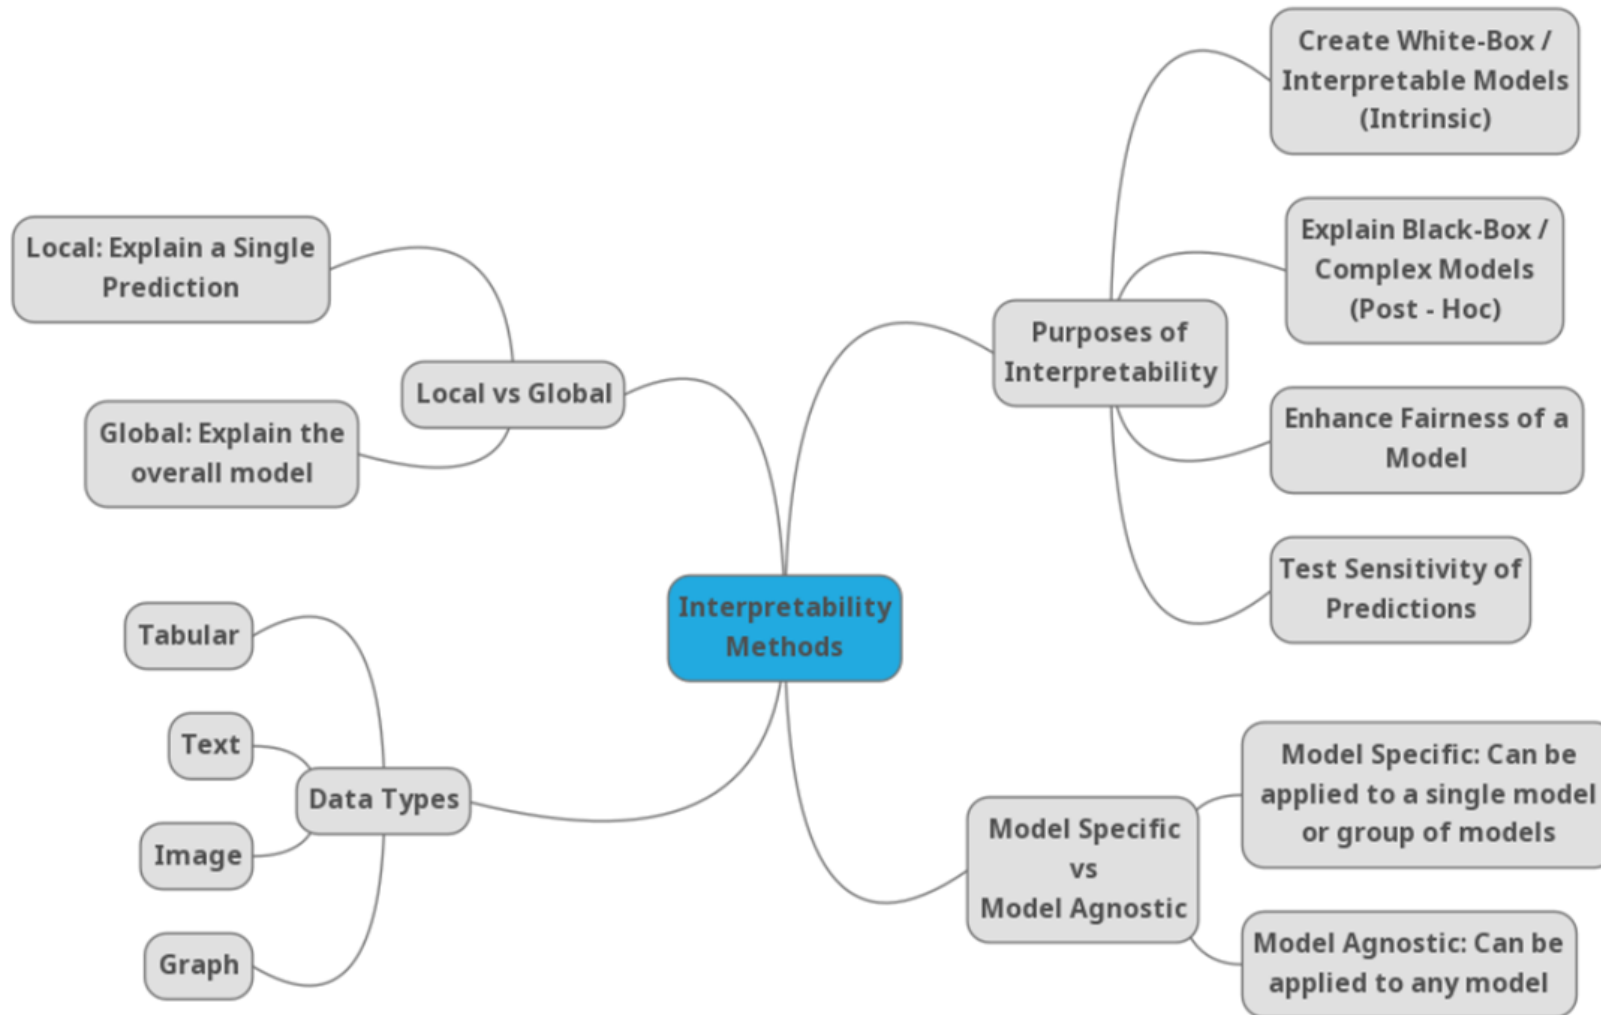


Figure 2. Taxonomy mind-map of Machine Learning Interpretability Techniques.

# Limits to ‘transparency’ in Explainable AI

*“Burrell (2016) distinguishes between three barriers to transparency:*

*(1) **intentional concealment on the part of corporations or other institutions**, where decision-making procedures are kept from public scrutiny;*

*(2) **gaps in technical literacy**, which mean that, for most people, simply having access to underlying code is insufficient; and*

*(3) a “**mismatch between the mathematical optimization in high-dimensionality characteristic of machine learning and the demands of human-scale reasoning and styles of interpretation.**” ...” - Goodman and Flaxman (2017)*

Goodman, B., & Flaxman, S. (2017). **European Union regulations on algorithmic decision-making and a “right to explanation”**. *AI magazine*, 38(3), 50-57.

Burrell, J. 2016. **How the Machine “Thinks”**: Understanding Opacity in Machine Learning Algorithms. *Big Data and Society* 3(1)

# Rational Extraction

*“ ... we learn to extract pieces of input text as justifications – rationales – that are tailored to be short and coherent, yet sufficient for making the same prediction.*

*... In order for the subset to qualify as a rationale it should satisfy two criteria:*

*1) the selected words should be interpretable and*

*2) they ought to suffice to reach nearly the same prediction (target vector) as the original input. ...”*

<b>Review</b>		
the beer was n't what i expected, and i'm not sure it's "true to style", but i thought it was delicious. <b>a very pleasant ruby red-amber color</b> with a relatively brilliant finish, but a limited amount of carbonation, from the look of it. aroma is what i think an amber ale should be - a nice blend of caramel and happiness bound together.		
<b>Ratings</b>		
	<b>Look: 5 stars</b>	Smell: 4 stars

**Figure 1:** An example of a review with ranking in two categories. The rationale for Look prediction is shown in bold.



# Data Analysis

## - Trustworthy AI -

# Definition of Trustworthy AI

## ETHICS GUIDELINES FOR TRUSTWORTHY AI: High-Level Expert Group on Artificial Intelligence

“Trustworthy AI has three components ...:

1. it should be **lawful**, complying with all applicable laws and regulations ;
2. it should be **ethical**, ensuring adherence to ethical principles and values; and
3. it should be **robust**, both from a technical and social perspective, since, even with good intentions, AI systems can cause unintentional harm.”

“In an AI context, equality entails that the system’s operations cannot generate unfairly biased outputs (e.g. the **data used to train AI systems should be as inclusive as possible, representing different population groups**).”

# Bias in Data

he (128)



tinkering (0.20)  
productive (0.19)  
tinker (0.18)  
play (0.17)  
mechanics (0.16)  
practice (0.14)  
improvise (0.14)  
hopefully (0.13)  
execute (0.13)

she (72)



volunteering (0.18)  
volunteer (0.18)  
caring (0.15)  
interns (0.15)  
rehearse (0.13)  
rehearsals (0.13)  
professions (0.11)  
workers (0.11)  
labor (0.10)

# Data Statements

## Data Statements for Natural Language Processing: Toward Mitigating System Bias and Enabling Better Science

**Emily M. Bender**  
Department of Linguistics  
University of Washington  
ebender@uw.edu

**Batya Friedman**  
The Information School  
University of Washington  
batya@uw.edu

“A **data statement** is a characterization of a dataset that provides context to allow developers and users to better understand ... **what biases might be reflected in systems built on the software.**”

“We propose here that **foregrounding the characteristics of our datasets** can help, by allowing reasoning about what the **likely effects** may be and by making it clearer **which populations are and are not represented**, for both training and test data.”





# Data Justice



INTERNET POLICY REVIEW  
Journal on internet regulation

OPEN ACCESS

PUBLISH

ARCHIVES

ABOUT



Diversity

Governance

Infrastructure & Standards

Information & Data

Innovation

Intellectual Property Rights

Privacy & Security



Volume 11, Issue 1 | Concepts of the digital society

## Data justice



Lina Dencik, *School of Journalism, Media and Culture, Cardiff University, United Kingdom, dencikl@cardiff.ac.uk*  
Javier Sanchez-Monedero, *University of Córdoba, Spain*

PUBLISHED ON: 14 Jan 2022 DOI: 10.14763/2022.1.1615



### ABSTRACT

Data justice has emerged as a key framework for engaging with the intersection of datafication and society in a way that privileges an explicit concern with social justice. Engaging with justice concerns in the analysis of information and communication systems is not in itself new, but the concept of data justice has been used to denote a shift in understanding of what is at stake with datafication beyond digital rights. In this essay, we trace the lineage and outline some of the different traditions and approaches through which the concept is currently finding expression. We argue that in doing so, we are confronted with tensions that denote a politics of data justice both in terms of what is at stake with datafication and what might be suitable responses.

PDF

CITE

TWEET



DYSLEXIA MODE

Adjusts contrasts, text, and spacing in order to improve legibility for people with dyslexia.

### FEEDBACK:

Is this feature helpful for you, or could the design be improved? If you have feedback please send us a message.

### METRICS



OLLSCOIL NA GAILLIMHE  
UNIVERSITY OF GALWAY

Dencik, L., & Sanchez-Monedero, J. (2021). Data justice. *Internet Policy Review*, 11(1), 1-16.

# Linguistic Justice

*“... equitable access to social, economic, and political life regardless of linguistic repertoire”*

*“... linguistically just **NLP tools** must:*

*(1) work well for users regardless of language variety they use and*

*(2) work to counteract inequities based on language use in decision-making and resource allocation.*

*To build **linguistically just NLP tools**, we must recognize and address power inequities such as **over/underrepresentation of linguistic patterns** and discourses within datasets.”*



# Bias in NLP Data Sets

Task	Example of Representation Bias in the Context of Gender	D	S	R	U
Machine Translation	Translating “He is a nurse. She is a doctor.” to Hungarian and back to English results in “She is a nurse. He is a doctor.” (Douglas, 2017)		✓	✓	
Caption Generation	An image captioning model incorrectly predicts the agent to be male because there is a computer nearby (Burns et al., 2018).		✓	✓	
Speech Recognition	Automatic speech detection works better with male voices than female voices (Tatman, 2017).			✓	✓
Sentiment Analysis	Sentiment Analysis Systems rank sentences containing female noun phrases to be indicative of anger more often than sentences containing male noun phrases (Park et al., 2018).		✓		
Language Model	“He is doctor” has a higher conditional likelihood than “She is doctor” (Lu et al., 2018).		✓	✓	✓
Word Embedding	Analogies such as “man : woman :: computer programmer : homemaker” are automatically generated by models trained on biased word embeddings (Bolukbasi et al., 2016).	✓	✓	✓	✓

Table 1: Following the talk by Crawford (2017), we categorize representation bias in NLP tasks into the following four categories: (D)enigration, (S)tereotyping, (R)ecognition, (U)nder-representation.

# Data Detoxing

## Leashing the Inner Demons: Self-Detoxification for Language Models

Canwen Xu, Zexue He, Zhankui He, Julian McAuley  
University of California, San Diego  
{cxu, zehe, zhh004, jmcauley}@ucsd.edu

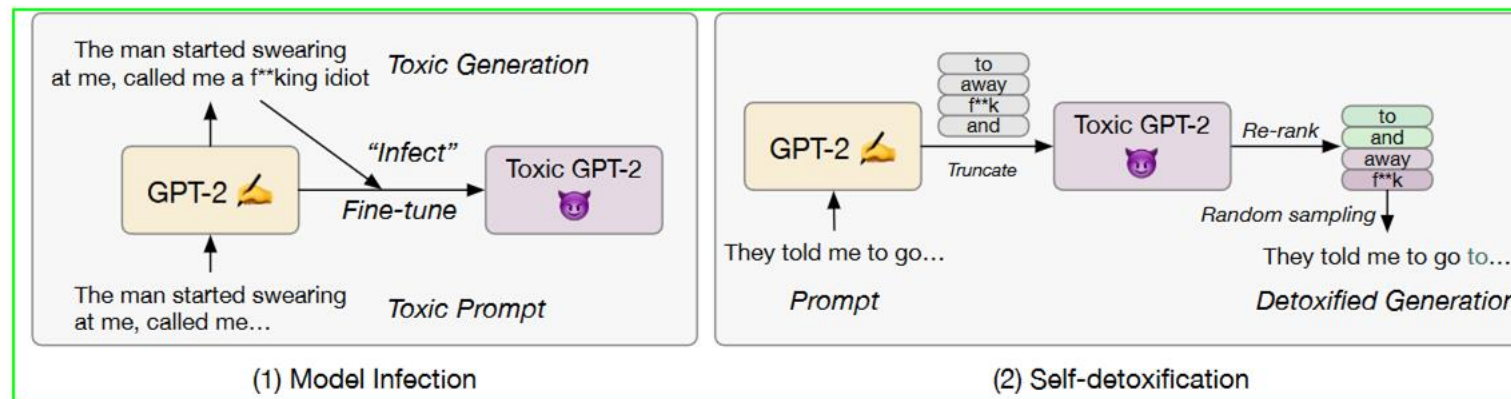


Figure 1: The workflow of self-detoxification. (1) We feed toxic prompts to the pretrained GPT-2 model to encourage toxic content to be generated. Then, we fine-tune a GPT-2 model on the generated toxic content and obtain an “infected” toxic GPT-2. (2) When doing self-toxification, the original GPT-2 model generates a probability distribution for the next token. After applying top- $k$  truncation, we use the toxic GPT-2 to score the token candidates and re-rank. Therefore, the words that are less favored by the toxic GPT-2 would have a better chance to be generated.

# AI for Good

## - NLP in Ethical Use Cases -



08-14-20

# Facebook's AI for detecting hate speech is facing its biggest challenge yet

Advancements in AI have dramatically improved the company's ability to identify written hate speech. But when it comes to rooting out hateful images, videos, and memes, Facebook's AI has a long way to go.

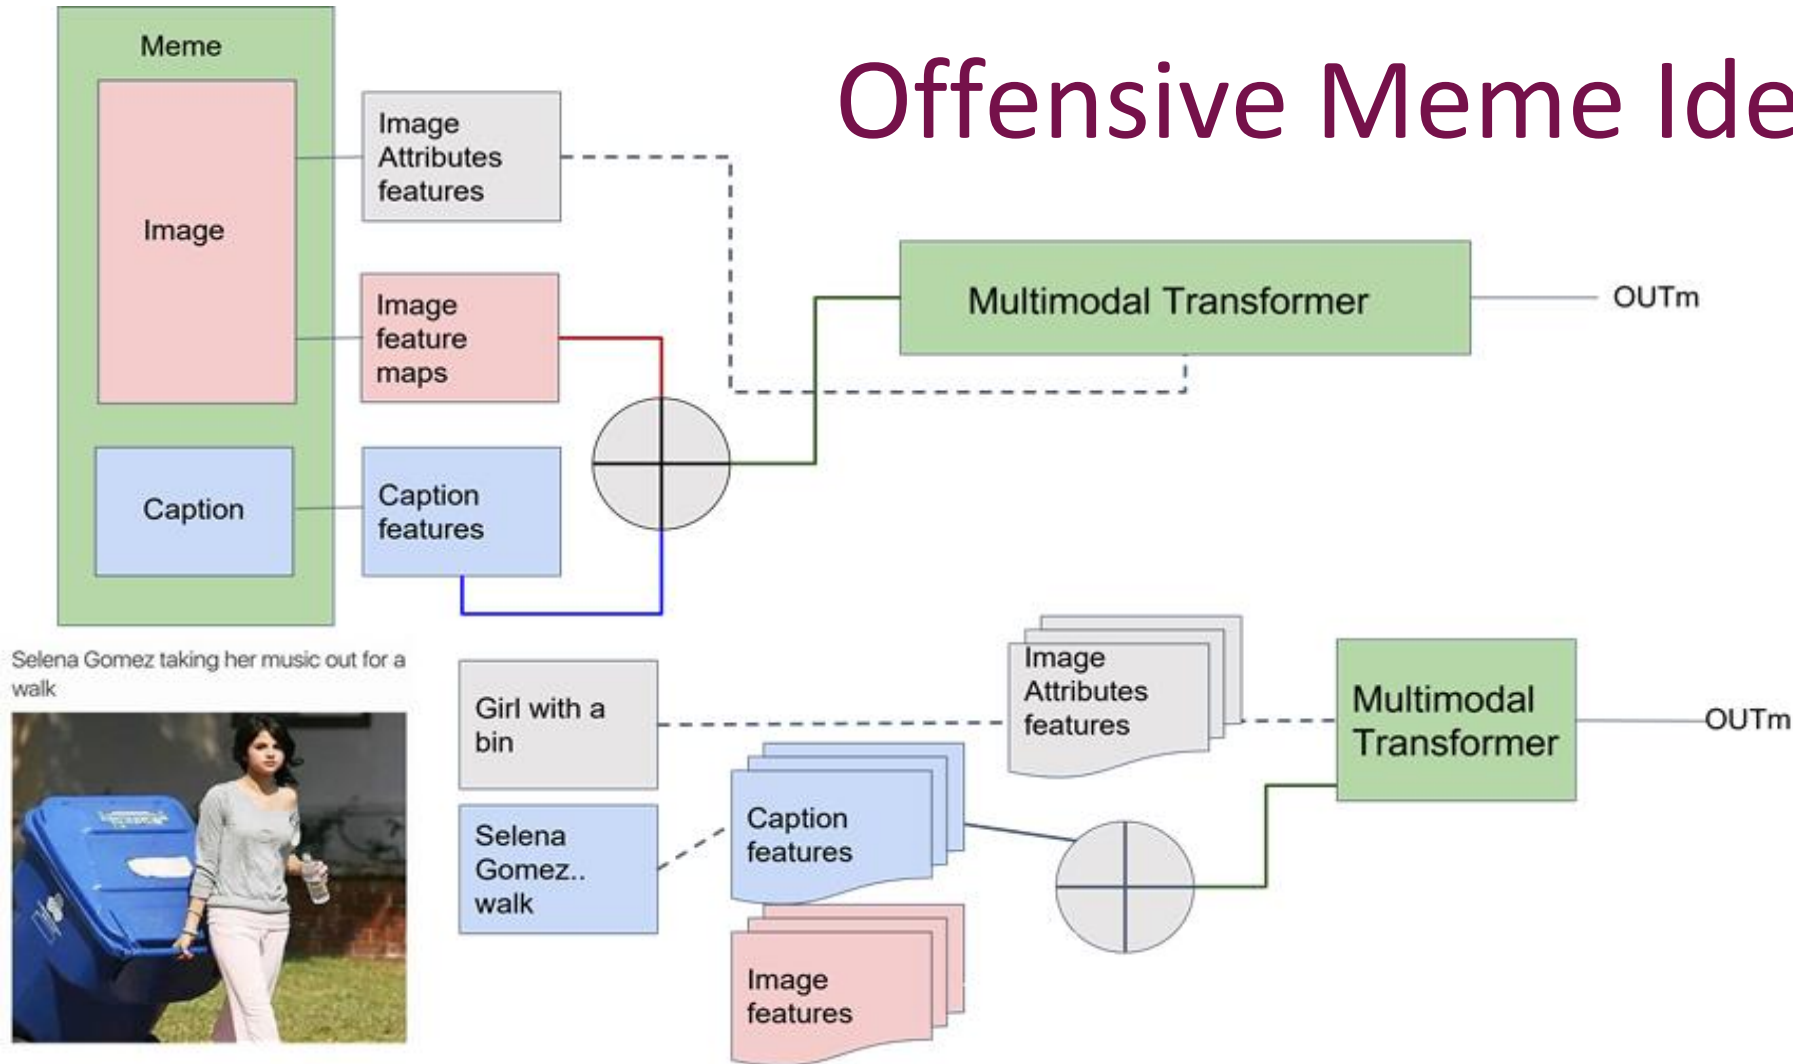
## Technology Networks Vaccines

[Find the science...](#)[Search](#)[Communities](#) ▾[Content](#) ▾[Webinars](#)[Home](#) > [Vaccines](#) > [News](#) > [Content Piece](#)

## Twitter Study Offers Insights Into Social Media's Influence on Vaccination

Published: May 16, 2022 | [Original story from New York University](#)

# Offensive Meme Identification



Insight 

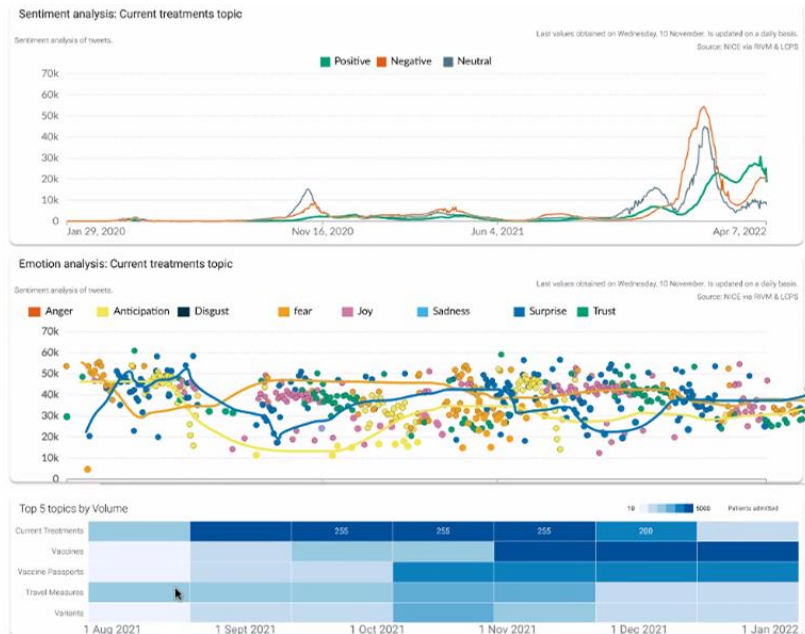
Suryawanshi, S., Chakravarthi, B. R., Arcan, M., & Buitelaar, P. (2020). **Multimodal meme dataset (MultiOFF) for identifying offensive content in image and text**. In *Proceedings of the Second Workshop on Trolling, Aggression and Cyberbullying* (pp. 32-41).

# Data Ethics Challenges

Bias in training data, resulting in **biased hate speech predictions**

- “... **dialect can lead to racial bias** in automatic hate speech detection models ... unexpected correlations between surface markers of African American English ...” – Sap et al 2019
- “... **impact of political bias** on hate speech classification by constructing three politically-biased data sets (left-wing, right-wing, politically neutral) ... political bias negatively impairs the performance of hate speech classifiers ...” – Wich et al 2020

**Social media analysis (emotion/sentiment, suggestions) to identify ‘two-way’ communication aspects around pandemics, such as reactions to government measures**



**Jason Y. Ng** @jasonyng · Apr 1

When Hong Kong announced a 14-day shutdown of leisure venues last week, it did a **half-baked job** banning only cinemas, gyms and a few other categories. It took a mini-community outbreak at a karaoke joint for the gov't to extend the list to include karaoke and mahjong parlors.

**Billie Dollins** @BillieDollins

Replying to @marcorubio and @POTUS

It would help if TEACHERS and school staff were vaccinated!!



# Data Ethics Challenges

- *“... availability, quality and nature of the training data ...*
- *... ability to **de-bias data** ... specific age and socioeconomic groups ... data from Facebook is likely to be **biased towards health data and linguistic quirks specific to a population** older than one trained on data from Snapchat ...*
- *... **public perception of privacy and data access** ... recent **survey of social media users** found that the majority considered analysis of their social media data to identify mental health issues “intrusive and exposing” and they **would not consent to this** ...*
- *... **assessment and evaluation of NLP models** to ensure that they are working as intended ... important **not to equate high scores with true language understanding** ...”*



# Learning Outcomes

After completing this topic, you should be able to:

- understand the challenge of ethical issues in NLP
- be aware of major regulatory frameworks that address Ethics in AI
- understand some of the main concepts in research on ethical issues in NLP

# Lab of this week

Exercises on bias in NLP



# Industry Talk - Genesys



**Dr. Maciej Dąbrowski**

Chief Data Scientist, Digital & AI  
Genesys, Galway

