



# Principles of Machine Learning


## Week 8: Linear Regression in One and Multiple Variables (Part 1)



# Learning Objectives

After successfully completing this topic, you will be able to ...

- Explain what Linear Regression is and its use in ML
- Describe and implement a Gradient Descent algorithm for learning LR parameters with one or multiple variables
- Describe and implement extensions such as Stochastic GD, regularisation, and polynomial regression
- Considering the characteristics of linear regression, recommend when it would be appropriate to an application
- Discuss and apply feature engineering methods including feature scaling, feature reduction, and transformation methods (e.g. PCA)



# Structure of Videos for This Topic

## Week 8

- Linear regression; closed-form solution for linear regression
- Gradient descent for linear regression with one input variable
- Multiple Linear Regression: solving in closed form and with gradient descent

## Week 9

- Feature scaling; polynomial regression
- Bias, variance, underfitting and overfitting in regression
- Gradient descent with regularisation; dimension reduction
- Feature engineering approaches including transformations and subset selection





# Part 1: Linear Regression



# Linear Regression Overview

- Given target value ( $y$ ) and set of attribute values ( $x_m$ ), goal of Linear Regression is to find an equation that describes the target in terms of the attributes


- Equation is of the form:


$$y = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_m x_m$$

where  $\theta_m$  is the weight associated with attribute  $x_m$

- This is the equation of a line

- Compare to a 2D line, where  $x$  and  $y$  are dimensions:  $y = a + b x$
- Now have higher dimensions, instead of  $x$  we have  $x_1, x_2$ , etc
- Use  $\theta_0$  as intercept rather than  $a$ , and  $\theta_1$  relates to slope in dimension  $x_1$





# Linear Regression Overview

- Have to find equation of this form:

$$y = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_m x_m$$

where  $\theta_m$  is the weight associated with attribute  $x_m$

- In other words, have to find values of  $\theta_0 \theta_1 \theta_2$  etc
- Weights found using least squares fit
  - For a given training set, the weights are found that minimize the squared error between predicted and actual target values
  - If a weight is 0, that attribute has no effect on outcome
  - Assumes that there is little/no correlation between dimensions





# Linear Regression Overview

- Training data consists of sets of values for

$x_1 \ x_2 \dots y$

Size (m <sup>2</sup> )	# Beds	# Floors	Age (yrs)	Price (k€)
195	5	1	40	450
130	3	2	35	220
140	3	2	26	310
80	2	1	30	170
180	5	2	38	400

$x_1 \quad x_2 \quad x_3 \quad x_4 \quad y$

- We formulate a hypothesis of the form


$$h_{\theta}(\mathbf{x}) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_m x_m$$

where  $\mathbf{x}$  is vector  $[x_1 \ x_2 \dots]$  and  $\theta$  is vector  $[\theta_0 \ \theta_1 \ \theta_2 \dots]$

- **Learning objective:** find values for  $\theta$  such that the value output by  $h_{\theta}(\mathbf{x})$  for a given vector  $\mathbf{x}$  is as close to  $y$  as possible



# Linear Regression in 1 Variable



Training data: green +

Hypothesis: blue line.

Hypothesis  $h_{\theta}(x)$


has the form  $h_{\theta}(x) = \theta_0 + \theta_1 x$

$\theta_0, \theta_1$  are the **model parameters/weights**.





# Linear Regression in 1 Variable



Training data: green +

Hypothesis: blue line.

Hypothesis  $h_{\theta}(x)$   
has the form  $h_{\theta}(x) = \theta_0 + \theta_1 x$

$\theta_0, \theta_1$  are the **model parameters/weights**.

**Learning objective:** choose  $\theta_0, \theta_1$  such that for each training example  $(x^{(i)}, y^{(i)})$ ,  $h_{\theta}(x^{(i)})$  is as **close as possible** to  $y^{(i)}$  on average.

One way to formalise this is mean squared error:

$$\min_{\theta_0, \theta_1} \frac{1}{2N} \sum_{i=1}^N (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

This defines a **squared error cost function**,  $J(\theta_0, \theta_1)$ :

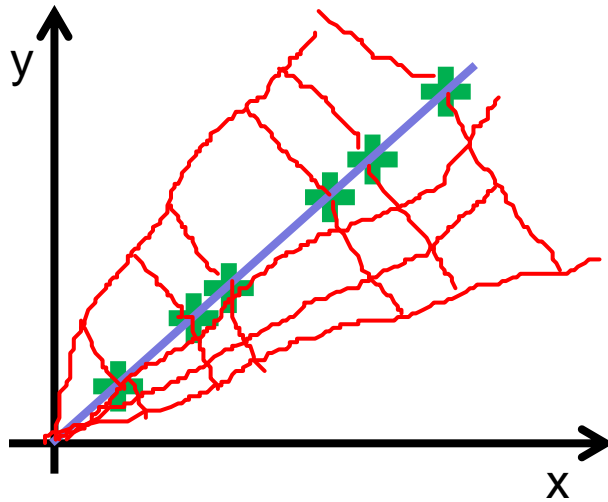
$$J(\theta_0, \theta_1) = \frac{1}{2N} \sum_{i=1}^N (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

Minimize  $J$  to find the optimal hypothesis.

Note: could use a different cost function; may yield slightly different results. Squared loss dates to Gauss.




# Cost Function Behaviour [1]



Simplified case:  
data values are  $y=x$


Hypothesis form is  $h_{\theta}(x) = \theta_1 x$   
 $\theta_0 = 0$  (intercept is 0)  
Only vary  $\theta_1$




**If** there is a straight line that goes through the data, it is defined by the parameters of  $\theta$  where  $J(\theta) = 0$ , irrespective of form of  $J$ .

**If not**, min value of  $J$  will be positive and will depend on specific cost fn used.

**Either way**,  $J$  is convex, has single minimum.




## Cost Function Behaviour [2]



If we need to optimize  $\theta_0$  and  $\theta_1$ :  
still have single minimum, still a convex optimization problem.

This also holds when we move to linear regression in multiple variables, with more  $\theta_i$ .





# How Do We Optimize the Parameters?

There are two main ways to find  $\theta_0$  and  $\theta_1$ :

1. Closed form solution:

- Because it is a convex optimization problem, the solution has a unique form

2. Gradient Descent:

- General-purpose algorithm for finding a local minimum of any continuous differentiable function
- For convex hull, it can always find the correct answer if its parameters are reasonable
- Iterative unlike closed form and therefore, not as efficient as closed form
- Advantage: Applicable to many optimisation tasks

Either way, need partial derivatives of  $J(\theta_0, \theta_1) \dots$



# Cost Function Derivatives

Given a cost function  $J(\theta_0, \theta_1 \dots \theta_m)$ ,  
its minimum value w.r.t. all  $\theta_i$  is found when all of its partial  
derivatives are zero.


In this case:

$$J(\theta_0, \theta_1) = \frac{1}{2N} \sum_{i=1}^N (h_{\theta}(x^{(i)}) - y^{(i)})^2 = \frac{1}{2N} \sum_{i=1}^N (\theta_0 + \theta_1 x^{(i)} - y^{(i)})^2$$

Partial derivatives:

$$\frac{\partial}{\partial \theta_0} J(\theta_0, \theta_1) = \frac{1}{N} \sum_{i=1}^N (\theta_0 + \theta_1 x^{(i)} - y^{(i)}) = 0$$

$$\frac{\partial}{\partial \theta_1} J(\theta_0, \theta_1) = \frac{1}{N} \sum_{i=1}^N (\theta_0 + \theta_1 x^{(i)} - y^{(i)}) x^{(i)} = 0$$



# Cost Function: Closed Form Solution

The partial derivatives have a unique solution:

$$\theta_0 = \frac{1}{N} \sum y^{(i)} - \frac{\theta_1}{N} \sum x^{(i)}$$

$$\theta_1 = \frac{N \sum x^{(i)} y^{(i)} - (\sum x^{(i)}) (\sum y^{(i)})}{N \sum \left( \underline{x^{(i)}}^2 \right) - (\sum \underline{x^{(i)}})^2}$$

This can be computed with a small amount of code or even in a spreadsheet.

[LinRegClosedForm.xlsx](#)





# Part 2: Gradient Descent



# Gradient Descent

- General-purpose method that works **beyond linear regression**: does not require closed-form solution
- Make initial guess; take incremental steps 'downhill' with step size controlled by **learning rate  $\alpha$**  until little/no change


Batch Gradient Descent **Algorithm**:

**initialise  $\theta$**  to any set of valid initial values


**repeat** until convergence (or time limit reached):

**simultaneously foreach  $\theta_j$  in  $\theta$  do:**


$$\theta_j \leftarrow \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$








# Gradient Descent Illustration



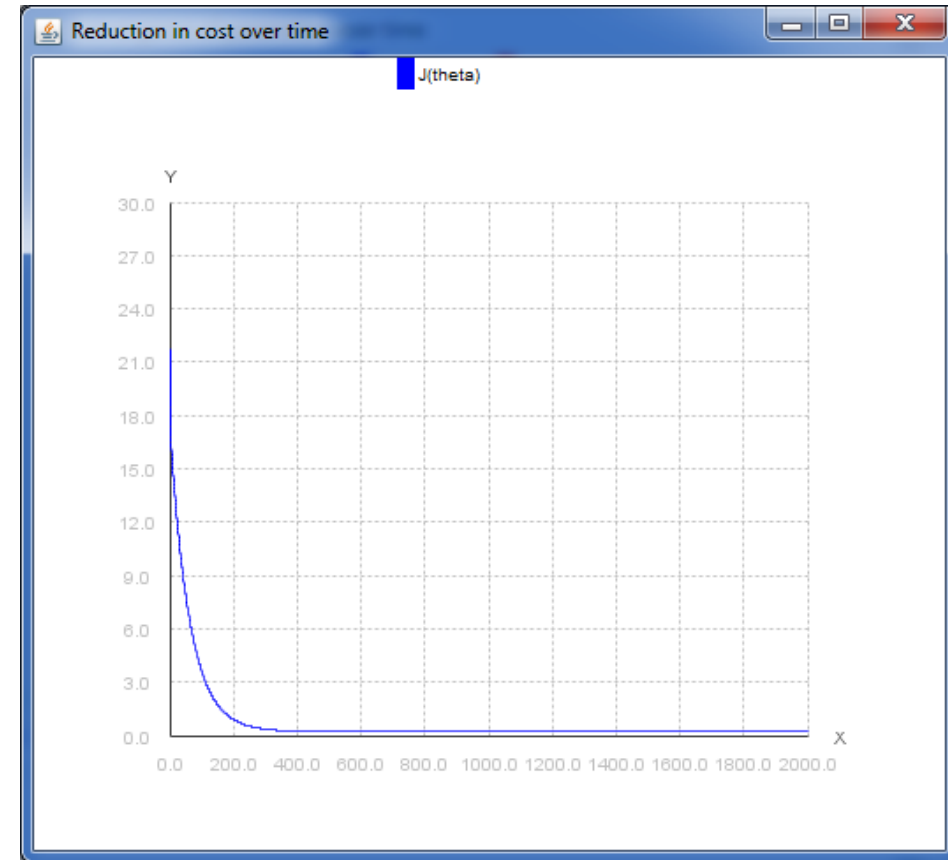




# Checking for Convergence



Changes to  $\theta_i$  will be initially large,  
reducing as iterations proceed.  
If progress is slow, increase learning rate.




Cost **must decrease** in every iteration,  
otherwise  $\alpha$  is too large



# Gradient Descent Illustration

- I have written two implementations of Gradient Descent for Linear Regression in one variable, both of which are on Blackboard
- Java version:  
`GradientDescent.java`
- Python (Jupyter Notebook) version:  
`SimpleLinearRegression.ipynb`
- We will examine the second one ...



# Stochastic Gradient Descent

- Rather than using all data together ("batch") in the update function, randomly select a single example for updating  $\theta$ 
  - Keep repeating until convergence
  - Partial derivatives are simplified: only 1 example so  $N=1$ , no  $\Sigma$

$$\frac{\partial}{\partial \theta_1} J(\theta) = \cancel{\frac{1}{N}} \sum_{i=1}^{\cancel{N}} (h_{\theta}(\mathbf{x}^{(i)}) - y^{(i)}) x_1^{(i)}$$






# Stochastic Gradient Descent

- Can be faster than batch gradient descent
  - Convergence is not guaranteed
- Works well in **online learning**
  - New data arriving all of the time
  - May be *high velocity* data
  - As long as all data are identically drawn from same distribution, just sample the data stream as fast as possible for updates



# Part 3: Multiple Linear Regression



# Multiple Linear Regression [1]

- Usually, have multiple input variables that relate to a quantity of interest

	Size (m <sup>2</sup> )	# Beds	# Floors	Age (yrs)	Price (k€)
	195	5	1	40	450
	130	3	2	35	220
$\mathbf{x}^{(3)}$	140	3	2	26	310
$y^{(3)}$	80	2	1	30	170
	180	5	2	38	400

$x_2^{(3)}$

$x_1$   $x_2$   $x_3$   $x_4$   $y$

N  
no. of cases

No. of attributes:  $m$

Now,  $\mathbf{x}^{(i)}$  is a feature vector of length  $m$ , values from one row

Individual features denoted  $x_j^{(i)}$





## Multiple Linear Regression [2]

Hypothesis is now of the form:

$$h_{\theta}(\mathbf{x}) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_m x_m$$

To simplify notation, add a dummy input variable  $x_0 = 1$ .

$$\text{Then: } h_{\theta}(\mathbf{x}^{(i)}) = \boldsymbol{\theta} \cdot \mathbf{x}^{(i)} = \sum \theta_j x_j^{(i)}$$

Cost function essentially same as before, now using vector notation:

$$J(\boldsymbol{\theta}) = \frac{1}{2N} \sum_{i=1}^N (h_{\theta}(\mathbf{x}^{(i)}) - y^{(i)})^2$$

Reminder of vector dot product:

$$\mathbf{a} = \langle a_1 \ a_2 \ a_3 \ a_4 \rangle$$

$$\mathbf{b} = \langle b_1 \ b_2 \ b_3 \ b_4 \rangle$$

$$\mathbf{a} \cdot \mathbf{b} = a_1 b_1 + a_2 b_2 + \dots$$



# Multiple LR: Closed Form Solution

As before, can be solved in closed form.

Let  $\mathbf{X}$  be full data matrix of inputs: row  $i$  corresponds to  $\mathbf{x}^{(i)}$

Let  $\mathbf{y}$  be vector of all outputs

Then the optimal vector of values for  $\boldsymbol{\theta}$  that minimises squared error is found from:

$$\boldsymbol{\theta}^* = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$$

Non-iterative, but requires inverting a matrix of size  $m \times m$   
( $m$  = no. of attributes):  $O(m^3)$

Gradient descent faster for large  $m$  (e.g.  $>1000$ ).





# Multiple LR: Gradient Descent

The algorithm is **the same** for Multiple LR as before:

**initialise**  $\theta$  to any set of valid initial values

**repeat** until convergence or until limit is reached:

**simultaneously foreach**  $\theta_j$  in  $\theta$  do:

$$\theta_j \leftarrow \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$

The partial derivatives are **different**: can be verified to be direct generalisation of the single-variable case:

$$\frac{\partial}{\partial \theta_j} J(\theta) = \frac{1}{N} \sum_{i=1}^N (h_{\theta}(\mathbf{x}^{(i)}) - y^{(i)}) x_j^{(i)}$$




# Multiple LR: Gradient Descent

- Let's compare the partial derivatives ...

- Single variable case: 
$$\frac{\partial}{\partial \theta_0} J(\theta_0, \theta_1) = \frac{1}{N} \sum_{i=1}^N (h_{\theta}(x^{(i)}) - y^{(i)})$$

$$\frac{\partial}{\partial \theta_1} J(\theta_0, \theta_1) = \frac{1}{N} \sum_{i=1}^N (h_{\theta}(x^{(i)}) - y^{(i)}) x^{(i)}$$

- Multiple variable case: 
$$\frac{\partial}{\partial \theta_j} J(\boldsymbol{\theta}) = \frac{1}{N} \sum_{i=1}^N (h_{\boldsymbol{\theta}}(\mathbf{x}^{(i)}) - y^{(i)}) x_j^{(i)}$$



*End of*

# Linear Regression in One and Multiple Variables (Part 1)